

Robust Platform for Scientific Computing: Python

Yung-Yu Chen

`yyc@solvcon.net`

Python Hsinchu User Group

September 24, 2013

Topics


- 1 The Need for Scientific Computing
 - The State of the Art
 - Python Supports Scientific Computing
- 2 The Python Ecosystem
 - Getting Start
 - Scientific Python Toolkits
- 3 HPC Code Development for Research
 - First-Principle Simulations and Conservation Laws
 - SOLVCON
 - Software System Structure
 - Parallel Computing

Supercomputing



- Specialized hardware for specific applications.
- **Speed** is the number 1 objective.
- Top 500 list (<http://top500.org/lists/2013/06/>):
 - 1 Tianhe-2: 3.12M cores, 33.86 Pflops (Peta floating-point operations per second). Equips Xeon E5-2692 and Xeon-Phi 31S1P.
 - 2 Titan (Cray XK7): 0.56M cores, 17.59 Pflops. Equips Opteron 6274 and NVIDIA K20x.
 - 3 Sequoia (IBM BlueGene/Q): 1.57M cores, 17.17 Pflops. Equips Power PQC 16C.

Cloud Computing



- Provide **elastic** computing power through virtualization technology.
 - Programs are run not on real hardware, but the virtualized systems.
 - Users can dynamically allocate resources including computing nodes and cores, memory, storage, etc.
- Pay-as-you-go allows everyone to solve significantly large problems.

But This Is Really We Are Using



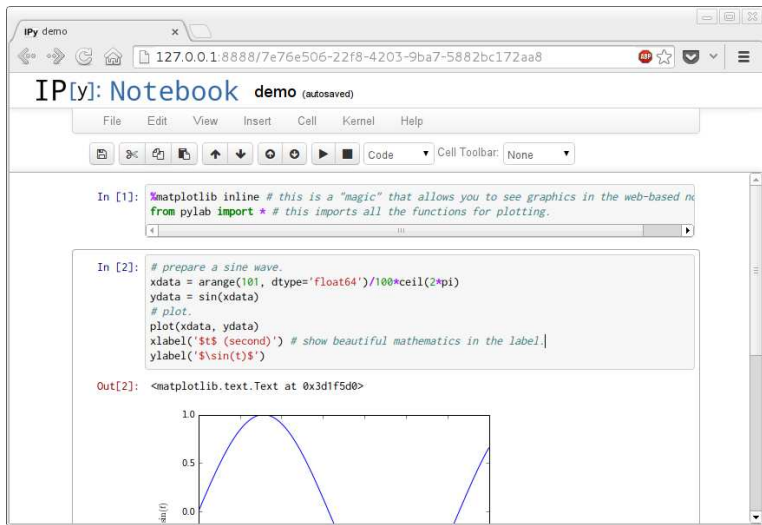
Programming Platform Matters

- You need a programming language, or several programming languages, to process, simulate, analyze, and exhibit the data of your problems.
 - It doesn't matter what supercomputers you are using.
 - It doesn't matter what cloud services you are using.
 - It doesn't matter if you are running the code on your laptop.
- You need to be proficient at a powerful programming language.
- And Python is that language.

What Python Can Do?

- An everyday tool:
`python -c 'import math; print math.factorial(10)'`.
- An established ecosystem for reproducible analysis.
 - A rich collection of scientific tools: NumPy, SciPy, Matplotlib, Pandas, iPython, etc.
- A popular platform for web programming.
 - Even include a web server:
`python -m SimpleHTTPServer.`
- A platform that allows you to do anything.
 - The Python Package Index (<https://pypi.python.org/pypi>) now has 34,924 packages.

A Rich Interactive Environment



Before Started

- The Python official site contains comprehensive information:
 - Table of contents of the documentation:
<http://docs.python.org/2/>.
 - The best reference to the standard library:
<http://docs.python.org/2/library/index.html>.
 - Read *The Python Tutorial*
(<http://docs.python.org/2/tutorial/index.html>).
- Start with Python 2.
- A good self-training material is *Learn Python the Hard Way*
(<http://learnpythonthehardway.org/book/>).

Installation

- Use Anaconda.
 - `https://store.continuum.io/cshop/anaconda/`.
 - It provides everything at once and can be easily updated.
- NOT recommended for scientific users:
 - Download from Python website.
 - Use your OS's package managers (apt-get, yum, ports, etc.)
- Why not?
 - They require expertise in Python and take long time.
 - They contains outdated packages.

*This suggestion is specific to scientific users.

Two Types of Programming

- Compiled vs interactive.
 - Low-level platform is built upon pre-compiled executables, like the Python runtime and the underneath libraries.
 - In a high-level interactive environment, we can type just several commands or press buttons to do complex computing.
- Python glues the low-level parts together, and expose them to the high-level.
- To configure and build the low-level platform needs a lot of efforts.
 - Products like Anaconda do that for us.

Categories of Scientific Python Tools

- Programming.
 - NumPy, Cython, iPython.
- Algorithms.
 - SciPy, SciKits.
- Visualization.
 - Matplotlib, VTK.
- Applications.
 - Pandas, NLTK, networkx, PyMOL, Pyomo, yt, ..., etc.

NumPy

- NumPy (<http://numpy.scipy.org/>) provides basic multi-dimensional array support.
- Array-oriented programming is the foundation to scientific computing.
- It provides basic facilities such as linear algebra and Fourier transform.

Let's see the demo.

Cython

- Cython (<http://cython.org/>) is a superset of the Python programming language.
- It speeds up Python code to be comparable of C.
- It provides interfaces for Python to use low-level C code or libraries.

Cython Benchmark

```

import numpy as np
cimport numpy as cnp

def action():
    cdef cnp.ndarray[cnp.double_t, ndim=2] arr0 = np.empty([1000,1000], dtype='float64')
    arr0.fill(0)
    cdef cnp.ndarray[cnp.double_t, ndim=2] arr1 = np.empty([1000,1000], dtype='float64')
    arr1.fill(1)
    cdef int it = 1
    cdef int jt
    while it < 999:
        jt = 1
        while jt < 999:
            arr0[it, jt] += arr1[it-1, jt-1]
            arr0[it, jt] += arr1[it-1, jt ]
            arr0[it, jt] += arr1[it-1, jt+1]
            arr0[it, jt] += arr1[it , jt+1]
            arr0[it, jt] += arr1[it+1, jt+1]
            arr0[it, jt] += arr1[it+1, jt ]
            arr0[it, jt] += arr1[it+1, jt-1]
            arr0[it, jt] += arr1[it , jt-1]
            jt += 1
        it += 1
    assert 7968032 == arr0.sum()

```

Cython is **41** times faster than normal Python.

iPython Notebook

- iPython stands for interactive Python.
- It provides plain-text, GUI, and web interface.
- iPython notebook is its web interface.
 - Use `ipython notebook` to launch.
 - Useful for research notes, experimenting, and education.
 - Mathematical expressions are the first-class citizen.
 - You can store and share any ipython notebook, even online: <http://nbviewer.ipython.org/>.

Let's see more demo of it.

SciPy Library

- The term SciPy has many meanings:
 - The SciPy library
(<http://docs.scipy.org/doc/scipy/reference/>);
what I want to talk about here.
 - The SciPy ecosystem (<http://www.scipy.org/>);
everything about Python for sciences.
 - The SciPy conference
(<http://conference.scipy.org/>); in North America,
Europe, and India.
 - The SciPy community; those who use Python for
scientific research.

Let's see how it works in an ipython notebook.

It's about SOLVCON


A solver constructor.

- Perform first-principle simulations for physical processes governed by conservation laws.
 - Usually formulated as hyperbolic partial differential equations (PDEs).
- Written in **Python** and with the performance hot-spot accelerated by C (or CUDA).
- Address **high-performance computing (HPC)** by mesh-based, array-oriented programming.

See <http://solvcon.net/> for detail.

Conservation Laws Govern The World


Fluid mechanics, solid mechanics, electromagnetism, etc.




Supersonic flow.



Atmospheric flow.



Seismic waves.



Electromagnetic waves.

First-Order Hyperbolic PDEs

- The fore-mentioned problems share a common trait: Demanding time-accurate solutions of conservation laws.
- So this is what I want to solve:

$$\frac{\partial u_m}{\partial t} + \sum_{\mu=1}^3 \frac{\partial f_m^{(\mu)}(\mathbf{u})}{\partial x_\mu} = s_m(\mathbf{u})$$
$$\Rightarrow \boxed{\oint_{S(V)} \mathbf{h}_m(\mathbf{u}) \cdot d\mathbf{a} = \int_V s_m(\mathbf{u}) dv} \quad (1)$$
$$m = 1, \dots, M.$$

Challenges in Programming

Coding for first-principle simulators is difficult. Why?

- 1 Recall the math:

$$\frac{\partial u_m}{\partial t} + \sum_{\mu=1}^3 \frac{\partial f_m^{(\mu)}(\mathbf{u})}{\partial x_\mu} = s_m(\mathbf{u}) \quad (1)$$

- 2 Various approaches to meshing and the associated data structures.
- 3 Parallel programming for HPC.
- 4 Data management and result analysis.

The CESE Method

- The space-time Conservation Element and Solution Element (CESE) method, developed by Chang at NASA Glenn.
 - Directly solves generic hyperbolic PDEs (Eq. (1)).
- Enable pluggable multi-physics in SOLVCON.
 - Compressible flows: $\mathbf{u} = (\rho, \rho v_1, \rho v_2, \rho v_3, \rho e)^t$.
 - Stress waves in solids:
 $\mathbf{u} = (v_1, v_2, v_3, \sigma_{11}, \sigma_{22}, \sigma_{33}, \sigma_{23}, \sigma_{13}, \sigma_{12})^t$.
 - Electromagnetic waves: $\mathbf{u} = (E_1, E_2, E_3, B_1, B_2, B_3)^t$.
 - Acoustics, shallow-water, viscoelasticity, etc.

Chang (1995) Journal of Computational Physics 119(2):295–324


Chen (2011), Ph.D. Dissertation

What Is SOLVCON


- A Python-based software framework for constructing time-accurate solvers of conservation laws for any physical processes.
- SOLVCON uses the CESE method.
 - Unstructured meshes of mixed elements are used in two- or three-dimensional space.
 - Message-passing is built into the framework for parallel computing.

Application: Stress Wave in Solids

- Beryl: Anisotropic crystal of hexagonal symmetry.






Exact solution of group velocity



Simulated result

Yang et al. (2011) J. Vib. Acoust. 133(2): 021001

Application: Supersonic Flows



2D cases:


- Flow over a cylinder.
- Oblique shock by a ramp.
- Moving shock climbing a ramp.
- Moving shock diffraction by a step.
- Moving shock past dust layer.
- Reflection of oblique shock.
- Implosion.

3D cases:

- Sod's shock tube.
- Flow over sphere.
- Jet in cross flow.

Jet in Supersonic Cross Flow

66 million elements are used in the simulation.




Density

Runtime Benchmark

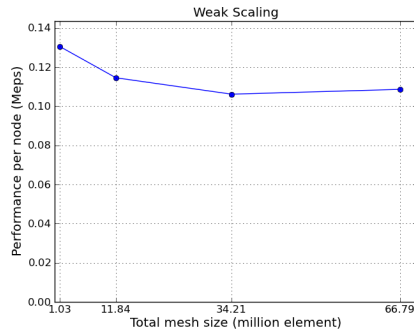
- Benchmark with hybrid parallel computing.
 - MPI across nodes; pthread within a node.
 - Run on Glenn@OSC: 4 cores/node with 10Gbps IB.
- Performance in million elements per second (Meps).

Number of cells (M)		1	11	34	66
Perf. (Meps)	1 core	0.035	—	—	—
	4 cores	0.13	—	—	—
	16 cores	0.45	0.47	—	—
	32 cores	—	0.91	—	—
	44 cores	—	1.26	1.33	—
	80 cores	—	—	2.16	—
	136 cores	—	—	3.61	3.82
	264 cores	—	—	—	7.17
	512 cores	—	—	—	12.7
	1024 cores	—	—	—	20.0

Scaling



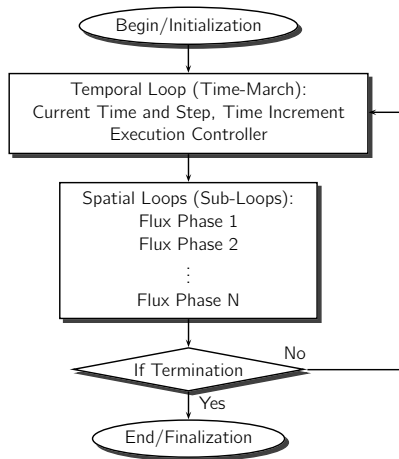
Fix Overall Mesh Size



Fix Per-Node Mesh Size


Two-Loop Structure of PDE Solvers

- The basic execution flow of SOLVCON:
 - Temporal loop for temporal (or pseudo-temporal) integration.
 - Spatial loops iterate over elements.
- The structure is general to all PDE solvers.



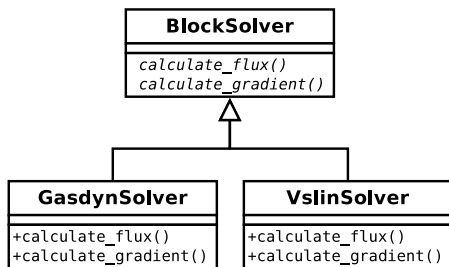
Solver Kernel for Spatial Loops

- A solver kernel is a Python class.
- The base class implements utility methods for **spatial loops**.
- The algorithms directly work with the mesh look-up tables.
- The concrete solver implements real algorithms, in **C**, or other fast languages.




Inheritance for Multi-Physics

- For a multi-physics algorithm, like the CESE method, a class hierarchy can be designed to host multiple physical processes.
- The physical processes are segregated.



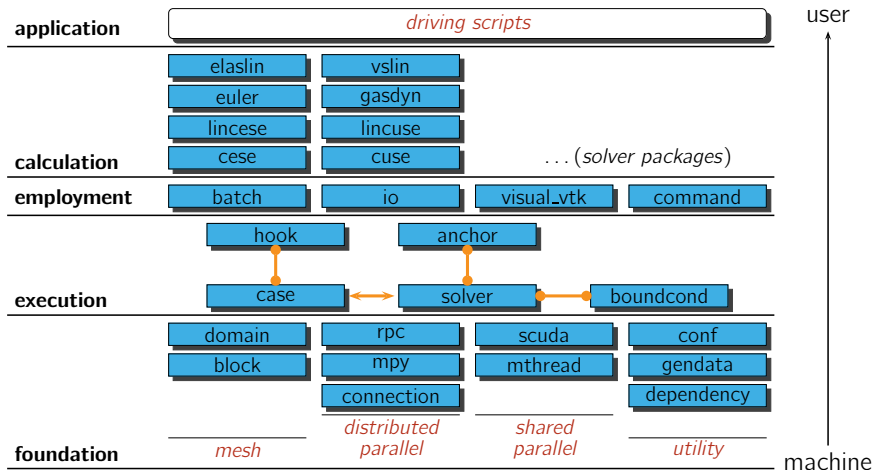
Temporal Loop and Call-Back

- A standalone class hierarchy (Case) is designed to host the **temporal loop**.

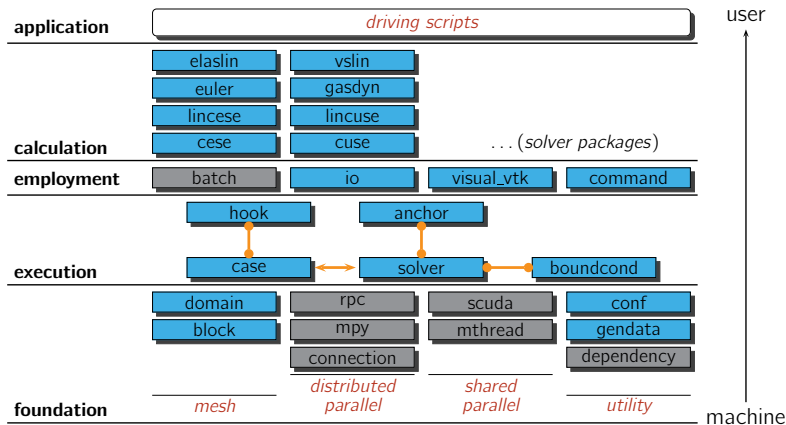


- Hook and Anchor are call-back objects for Case and Solver, respectively.
 - Supplement of main algorithms.
 - Lazy initialization.
 - Facilitating parallel computing and in-situ analysis.

Overall Design of SOLVCON




Renovation under Construction



- Simplify the architecture: Rely more on mpi4py, OpenMP, etc.
- Use Cython instead of ctypes for maintainability.


Two Types of Parallel Computing



- Simultaneously use shared-memory and distributed-memory parallel computing (DMPC & DMPC, respectively).
 - Main difference: **Addressing space**.
- Inter-process communication is needed.
 - DMPC is much more complex than SMPC.
 - **DMPC determines the scalability**.
 - **MapReduce is unsuitable**.


Extending Solver Kernel for SMPC

- A Solver class can be extended to use shared-memory parallel computing.
- Only the spatial loops are modified.
- Can use pthread, OpenMP, CUDA, OpenCL, etc.



Domain Decomposition for DMPC

- Before computation: Domain decomposition.
 - Use connectivity data to build the graph of cells.
 - Partition the graph by calling SCOTCH library.
 - Use the partitioned graph to decompose mesh data.

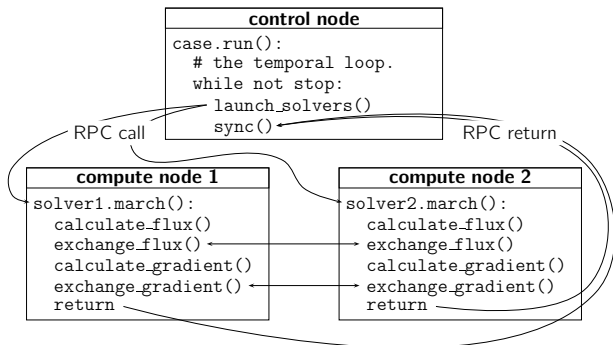


- During computation: Exchange data of the cells on the interface of different sub-domains.
 - Use MPI to communicate among sub-domains.

Solver Kernels Need Not Know DMPC

- DMPC is in SOLVCON framework.
- SMPC is in solver kernels.

DMPC Execution Flow in SOLVCON



- When developing solver kernels, we do not need to worry about the complexity of DMPC.
- Hybrid parallelism is achieved by the segregation.

Post-Processing is Bottleneck

- High-resolution simulations generate a lot of data:
 - For 50 million element mesh, the data for one scalar (single-precision) are **200 MB**.
 - A typical run has at least 10,000 time steps.
 - Transient analysis: $10,000 \times 200 \text{ MB} = \mathbf{2 \text{ TB}}$.
 - $\rho, p, T, \vec{v}, \vec{\omega}$ for CFD: $2 \text{ TB} \times 9 = \mathbf{18 \text{ TB}}$.
- Workaround: Reducing output frequency.
 - Every 100 time steps: **180 GB**.
- Post-processing the solutions is painfully time-consuming:
 - The large data are usually processed by using a single workstation.
 - Turnaround time could be in months.

Solutions in SOLVCON

- Parallel I/O.
 - Each sub-domain outputs its own solutions.
 - It is used with **parallel post-processing**.
- In situ visualization.
 - Visualization is being done on the fly with the simulation.
 - Everything happens in memory.
 - Output only graphic files, which are much smaller than the full solution field.
- Parallel I/O and in situ visualization are complementary to each other.

Python: Rich Ecosystem for Scientists

- Robust fundamental tools: NumPy, SciPy, Matplotlib, iPython, Cython.
- Abundant applications: Pandas, NLTK, networkx, VTK, PyMOL, Pyomo, ..., etc.
- iPython notebook is an excellent workbench from prototyping to presentation.
- Cython can help us to get the speed of C.
- A “virtual lab” can be built upon Python, like what SOLVCON is approaching.

Coding HPC for Research

- Identifying the fundamental structure.
 - In SOLVCON it's the two-loop structure.
 - Enabled by the insights from the “domain experts”.
- Use Python from the beginning and to the end.
 - Prototype your system with Python.
 - Gradually replace performance hot spots with Cython or low-level C.
 - Try to stay away from C++ or Fortran as much as possible.
- It's very **productive**.
 - SOLVCON is multi-physics by its clear structure with hybrid parallelism.
 - Python rocks.

Thanks!