

C/C++ for Python's Use: You Wrap It

Yung-Yu Chen

`yyc@solvcon.net`

The Happy Folk

December 26, 2013

Python is Slow, Way Too Slow

- The rule of thumb: Python is **two orders of magnitude slower** than C.
 - 10 milliseconds in C: 1 second in Python.
 - 8 hours in C: a month in Python.
- It's not OK to spend a month on an over-night job.
 - But we just can't lose the productivity by Python.
 - So we want to connect Python with C (and/or C++).
 - Another case: Python is the thin wrapper of the underlying system.

Ways to Wrapper

- Cython for C and Boost.Python for C++.
- There are other options:
 - ctypes and/or cffi.
 - SWIG.
 - <http://www.ohloh.net/p/pygccxml>.
 - ... and many that we don't talk about.

In Conclusion ...

- If you don't need to deal with C++.
Congratulations. You can just use Cython without worrying anything.
- If unfortunately you are working with C++. Go Boost.Python.
- <http://docs.python.org/2/c-api/index.html>
 - That is, use Python C API to make the low-level code available/callable in the Python interpreter.
 - Nothing stops you from wrapping by your bare hands with the C API.
 - But it's just an unpleasant approach (and error-prone).

Let's Start with the Bad Boy

- C++ and Python both rely on their own OO systems.
 - Both systems are powerful but substantially different.
- Boost.Python maps everything between C++ and Python.
 - Functions, classes, exceptions, containers, iterators, etc.
 - It's *like* writing Python extension with C++.

Official Boost.Python Hello World

```
1 #include <boost/python.hpp>
2 char const* greet() {
3     return "hello, world";
4 }
5 BOOST_PYTHON_MODULE(hello_ext) {
6     using namespace boost::python;
7     def("greet", greet);
8 }
```

- `#include <boost/python.hpp>`: turn on Boost.Python.
- `BOOST_PYTHON_MODULE`: make a Python module.
- `def("greet", greet)`: wrap a function.

Done.

Boost.Python Thinks from Bottom up

- It is a C++ library (part of **boost**). It operates around a fundamental C++ code base.
- To use the comprehensive wrapping capabilities, one has to write the code himself.
 - No code generation other than C++ template.
 - C++ only and only C++.

Exceptions

```
1  #include <boost/python.hpp>
2  #include <boost/python/errors.hpp>
3  void translate(const SomeException &err) {
4      PyErr(SetString(PyExc_MemoryError, err.what()));
5  }
6  BOOST_PYTHON_MODULE(error_ext) {
7      using namespace boost::python;
8      register_exception_translator<SomeException>(translate);
9  }
```

- `#include <boost/python/errors.hpp>`
- `translate()`
- `register_exception_translator`

Iterators

```

1  #include <boost/python.hpp>
2  #include <list>
3  struct Iterable {
4      typedef std::list<int>::iterator iterator;
5      iterator begin() { m_data.begin(); }
6      iterator end() { m_data.end(); }
7      std::list<int> m_data;
8  }
9  BOOST_PYTHON_MODULE(iter_ext) {
10     using namespace boost::python;
11     class_<Iterable>("Iterable", init<>())
12         .add_property("iter", range(Iterable::begin, Iterable::end))
13         .def("__iter__", range(Iterable::begin, Iterable::end))
14         ;
15 }

```

- `boost::python::range`: iterate from begin to end.
- Alternately, `__getitem__` and `__len__` can help.

Bi-Directional Conversions

- Python has GC while C++ doesn't: conversions of objects usually need to construct intermediate objects.
 - It's more complex when a container is involved.
- What to look for:
 - `to_python_converter`: take C++ object and expose it to Python.
 - `boost::python::converter::registry`: take Python object and give it to C++.
- When exchanging containers or large objects, using `boost::shared_ptr` as a holding type helps.

Obvious Caveats

- We need basic understanding about the CPython memory management.
 - Ownership of references for reference counting:
<http://docs.python.org/2/c-api/intro.html#reference-count-details>
- Be very careful about passing objects across the interpreter barrier.
 - Do not return intermediates, although Boost.Python checks for invalid returns.
 - Some return policy copies objects. It can hit the performance for large objects, or give wrong results for containers.

What Is Cython?

- Cython generates C code from Python.
 - We need to use its command-line (`cython`) like a compiler.
- The generated code contains Python C API calls.
 - Can interface to both the existing `C` and `C++` code.
- Some constructs can be translated into native C, and get utmost speed-up.
- <http://www.cython.org/>.

Cython is a Language

- Cython extends Python. It is a **superset** of Python.
 - When compiling **.py** files, of course only Python syntax can be used.
- If Cython-specific syntax is used, the source code can no longer be run by a vanilla Python VM.
 - It should be saved as a **.pyx** file and compiled by Cython compiler to binary.
- The Cython add-on helps to generate faster code and adapt to various libraries.
- It's a top-down tool, in contrast to Boost.Python.

Call C from Cython/Python

Let's say we have a C function `caction()`:

```

1  void caction(double arr0[][1000], double arr1[][1000]) {
2      for (int it=1; it<999; it++) {
3          for (int jt=1; jt<999; jt++) {
4              arr0[it][jt] += arr1[it-1][jt-1];
5              arr0[it][jt] += arr1[it-1][jt  ];
6              arr0[it][jt] += arr1[it-1][jt+1];
7              arr0[it][jt] += arr1[it  ][jt+1];
8              arr0[it][jt] += arr1[it+1][jt+1];
9              arr0[it][jt] += arr1[it+1][jt  ];
10             arr0[it][jt] += arr1[it+1][jt-1];
11             arr0[it][jt] += arr1[it  ][jt-1];
12         };
13     };
14 };

```

Cython Let You Call C

```
1  import numpy as np
2  cimport numpy as cnp
3  cdef extern:
4      void caction(double* arr0, double* arr1)
5  def action():
6      cdef cnp.ndarray[double, ndim=2] arr0 = np.empty(
7          [1000,1000], dtype='float64')
8      arr0.fill(0)
9      cdef cnp.ndarray[double, ndim=2] arr1 = np.empty(
10         [1000,1000], dtype='float64')
11      arr1.fill(1)
12      caction(&arr0[0,0], &arr1[0,0])
13      assert 7968032 == arr0.sum()
```

- `cdef extern:` tells Cython that we have a C function to be used.
- Then we just use it. Cython takes care of the Python C API boilerplates.

Useful Syntax for Wrapping

- `cdef extern from "header.h":`
 - Allows Cython to check the declarations from outside headers.
- `cdef public:`
 - Allows Cython to generate C header files for inclusion in other C code.

Cython Is for Number-Crunching

Consider the Python version "action":

```
1  import numpy as np
2  def action():
3      arr0 = np.empty([1000,1000], dtype='float64')
4      arr0.fill(0)
5      arr1 = np.empty([1000,1000], dtype='float64')
6      arr1.fill(1)
7      cdef int it = 1
8      cdef int jt
9      while it < 999:
10         jt = 1
11         while jt < 999:
12             arr0[it, jt] += arr1[it-1, jt-1]
13             arr0[it, jt] += arr1[it-1, jt ]
14             arr0[it, jt] += arr1[it-1, jt+1]
15             arr0[it, jt] += arr1[it , jt+1]
16             arr0[it, jt] += arr1[it+1, jt+1]
17             arr0[it, jt] += arr1[it+1, jt ]
18             arr0[it, jt] += arr1[it+1, jt-1]
19             arr0[it, jt] += arr1[it , jt-1]
20             jt += 1
21         it += 1
22     assert 7968032 == arr0.sum()
```

Let's See the Speed-Up

- C version runs 51.4 milliseconds; Python version runs 5.74 seconds.
- $5.74 \text{ sec} / 51.4 \text{ msec} = 106$ times faster.
- If our code is for number-crunching (calculation takes the majority of time), Python + Cython will deliver C-like performance.

Technique #1: Static Typing

- Just prefix “**cdef int**” to our counter:

```
def action():  
    cdef int it = 0  
    while it < 1000000:  
        it += 1  
    print it
```

- Then 16 times faster. (3.01 msec vs 49 msec)

```
> python -m timeit -s 'from example1d import  
    action' 'action()'  
100 loops, best of 3: 3.01 msec per loop
```

- Because Cython can use the type information to generate native C code instead of C API calls.

Technique #2: Direct Indexing

Provide array dimensions for direct indexing:

```
1  import numpy as np
2  cimport numpy as cnp
3  def action():
4      cdef cnp.ndarray[cnp.double_t, ndim=2] arr0 = np.empty([1000,1000], dtype='float64')
5      arr0.fill(0)
6      cdef cnp.ndarray[cnp.double_t, ndim=2] arr1 = np.empty([1000,1000], dtype='float64')
7      arr1.fill(1)
8      cdef int it = 1
9      cdef int jt
10     while it < 999:
11         jt = 1
12         while jt < 999:
13             arr0[it, jt] += arr1[it-1, jt-1]
14             arr0[it, jt] += arr1[it-1, jt ]
15             arr0[it, jt] += arr1[it-1, jt+1]
16             arr0[it, jt] += arr1[it , jt+1]
17             arr0[it, jt] += arr1[it+1, jt+1]
18             arr0[it, jt] += arr1[it+1, jt ]
19             arr0[it, jt] += arr1[it+1, jt-1]
20             arr0[it, jt] += arr1[it , jt-1]
21             jt += 1
22         it += 1
23     assert 7968032 == arr0.sum()
```

Direct Indexing Is Way Faster

- Use NumPy API for indexing: 5.74 sec.

```
> python -m timeit -s 'from example2d import  
    action' 'action()'   
10 loops, best of 3: 5.74 sec per loop
```

- Direct indexing: 139 msec.

```
> python -m timeit -s 'from example2a import  
    action' 'action()'   
10 loops, best of 3: 139 msec per loop
```

- Static typing and direct indexing make the code run **41** times faster.
- It's only 139 msec / 51.4 msec = 2.7 times slower than the C version.
 - Good enough for many uses.

Choose Your Wrapping Tool

- Python is two orders of magnitude slower than C.
- If you want to wrap C++, choose Boost.Python.
 - It thinks around C++. The wrapper should be built from bottom up.
- If you want to wrap C, choose Cython.
 - It thinks around Python. The wrapper should be built from top down.
- When working with NumPy arrays, static typing along with direct indexing makes the number-crunching very fast, just 2 times slower than the C counter part.