

# Implement High-Performance PDE Solvers for First-Principle Simulations by Using Python

Yung-Yu Chen  
yyc@solvcon.net

Developer  
SOLVCON Project

PyCon Japan 2012  
September 16

# It's about SOLVCON


- A solver constructor.
  - Perform first-principle simulations for physical processes governed by partial differential equations (PDEs).
  - I focus on conservation laws at the time being.
  - Written in **Python** and with the performance hot-spot accelerated by C (or CUDA).
  - Address **high-performance computing (HPC)** by mesh-based, array-oriented programming.
- Slides: [http://j.mp/yyc\\_pyconj2012](http://j.mp/yyc_pyconj2012)

# Outline

- 1 Background
  - First-Principle Simulations and Conservation Laws
  - Why Develop SOLVCON?
  - Applications of SOLVCON
- 2 The Design of SOLVCON
  - Mesh-Based Programming
  - Two-Loop Structure
  - Parallel Computing
  - Treating “Big Data”
- 3 Example: Sod’s Shock Tube
- 4 Conclusions

# Conservation Laws Govern The World


Fluid mechanics, solid mechanics, electromagnetism, etc.




Supersonic flow.



Atmospheric flow.



Seismic waves.



Electromagnetic waves.

# Experiments?



Hypersonic Wind Tunnel (Mach 5+)

# Experiments, Seriously?



# First-Order Hyperbolic PDEs

- The fore-mentioned problems share a common trait: Demanding time-accurate solutions of conservation laws.
- So this is what I want to solve:

$$\frac{\partial u_m}{\partial t} + \sum_{\mu=1}^3 \frac{\partial f_m^{(\mu)}(\mathbf{u})}{\partial x_\mu} = s_m(\mathbf{u})$$
$$\Rightarrow \boxed{\oint_{S(V)} \mathbf{h}_m(\mathbf{u}) \cdot d\mathbf{a} = \int_V s_m(\mathbf{u}) dv} \quad (1)$$
$$m = 1, \dots, M.$$

# Now a Programming Problem

Coding for first-principle simulators is difficult. Why?

- 1 Recall the math:

$$\frac{\partial u_m}{\partial t} + \sum_{\mu=1}^3 \frac{\partial f_m^{(\mu)}(\mathbf{u})}{\partial x_\mu} = s_m(\mathbf{u}) \quad (1)$$

- 2 Various approaches to meshing and the associated data structures.
- 3 Parallel programming for HPC.
- 4 Data management and result analysis.



# Why Write Code for Simulation?

Come on, there are many commercial codes. Is it really needed to develop yet another research code?

- Yes, because without the access to the source code, a computational scientist can never fully trust the code she uses.
- Commercial codes always implement the algorithms that can serve a larger crowd. We have to code the cutting-edge things by ourselves.

# Open-Source Research Codes?

- Good codes like FEniCS and FiPy didn't focus on time-accurate solution of conservation laws with HPC.
  - FEniCS is based on finite-element method (FEM).
  - FiPy is based on projection method.
- Plenty of research codes are open-source, but
  - As the name suggests, many research codes lack the modern structure that allows reuse.
  - 10 incompatible versions of a code in a research group? Nobody likes it, but it's quite often.

# And I Learned the CESE Method

- The space-time Conservation Element and Solution Element (CESE) method, developed by Chang at NASA Glenn.
  - Directly solves generic hyperbolic PDEs (Eq. (1)).
- Enable pluggable multi-physics in SOLVCON.
  - Compressible flows:  $\mathbf{u} = (\rho, \rho v_1, \rho v_2, \rho v_3, \rho e)^t$ .
  - Stress waves in solids:  
 $\mathbf{u} = (v_1, v_2, v_3, \sigma_{11}, \sigma_{22}, \sigma_{33}, \sigma_{23}, \sigma_{13}, \sigma_{12})^t$ .
  - Electromagnetic waves:  $\mathbf{u} = (E_1, E_2, E_3, B_1, B_2, B_3)^t$ .
  - Acoustics, shallow-water, viscoelasticity, etc.

Chang (1995) Journal of Computational Physics 119(2):295–324


Chen (2011), Ph.D. Dissertation

# Scope of SOLVCON


- A Python-based software framework for constructing time-accurate solvers of conservation laws for any physical processes.
- SOLVCON currently contains referential solvers that use the CESE method.
  - Unstructured meshes of mixed elements are used in two- or three-dimensional space.
  - Message-passing is built into the framework for parallel computing.
- A PDE-solving toolkit that can be embedded into other applications.

# Application: Stress Wave in Solids

- Beryl: Anisotropic crystal of hexagonal symmetry.






Exact solution of group velocity



Simulated result

Yang et al. (2011) J. Vib. Acoust. 133(2): 021001

# Application: Supersonic Flows



## 2D cases:


- Flow over a cylinder.
- Oblique shock by a ramp.
- Moving shock climbing a ramp.
- Moving shock diffraction by a step.
- Moving shock past dust layer.
- Reflection of oblique shock.
- Implosion.

## 3D cases:

- Sod's shock tube.
- Flow over sphere.
- Jet in cross flow.

# Jet in Supersonic Cross Flow

66 million elements are used in the simulation.



Density


# Runtime Benchmark

- Benchmark with hybrid parallel computing.
  - MPI across nodes; pthread within a node.
  - Run on Glenn@OSC: 4 cores/node with 10Gbps IB.
- Performance in million elements per second (Meps).


Number of cells (M)		1	11	34	66
Perf. (Meps)	1 core	0.035	—	—	—
	4 cores	0.13	—	—	—
	16 cores	0.45	0.47	—	—
	32 cores	—	0.91	—	—
	44 cores	—	1.26	1.33	—
	80 cores	—	—	2.16	—
	136 cores	—	—	3.61	3.82
	264 cores	—	—	—	7.17
	512 cores	—	—	—	12.7
	1024 cores	—	—	—	20.0



# Scaling




Fix Overall Mesh Size




Fix Per-Node Mesh Size

# Mesh-Based Programming

Curvilinear



Cartesian




Unstructured




- Two ways to discretize space:
  - Structured mesh: The connectivity of elements is structured. Cartesian is a special case.
  - Unstructured mesh: The connectivity of elements is **irregular**.
- SOLVCON uses **unstructured mesh**:
  - The data structure defines **connectivity** and **geometry**.
  - Code is based on the data structure.

# Unstructured Meshes of Mixed Elements

- Two-dimensional elements:




triangle




quadrilateral


- Three-dimensional elements:




tetrahedron



hexahedron




prism



pyramid

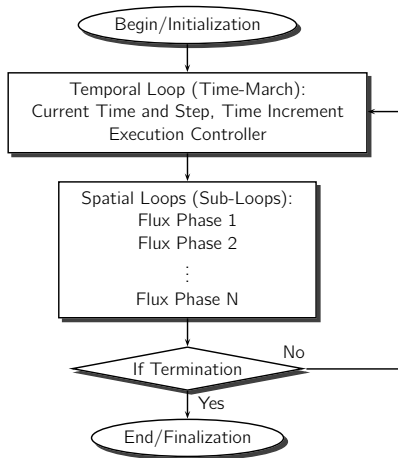
# Look-Up Tables for Meshes



- Consider cell-based solvers (value is stored at cell centers):
  - 4 nodes for the cell: 0, 1, 2, 3.
  - 4 faces for the cell: 012, 013, 123, 203.
- Tables to connect all cells:
  - Global index of nodes (ndcrd).
  - Nodes in each cell (c1nds).
  - Faces in each cell (c1fcs).
  - Nodes in each face (fcnds).
  - Cells connected by each face (fccls).


# Two-Loop Structure of PDE Solvers

- The basic execution flow of SOLVCON:
  - Temporal loop for temporal (or pseudo-temporal) integration.
  - Spatial loops iterate over elements.
- The structure is general to all PDE solvers.



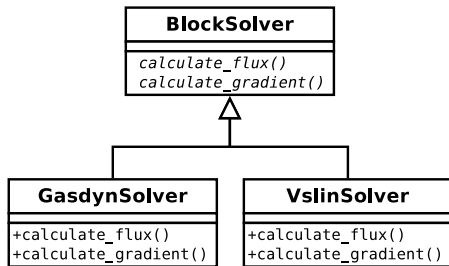
# Solver Kernel for Spatial Loops

- A solver kernel is a Python class.
- The base class implements utility methods for **spatial loops**.
- The algorithms directly work with the mesh look-up tables.
- The concrete solver implements real algorithms, in **C**, or other fast languages.



# Inheritance for Multi-Physics


- For a multi-physics algorithm, like the CESE method, a class hierarchy can be designed to host multiple physical processes.



- The physical processes are segregated.

# Temporal Loop and Call-Back

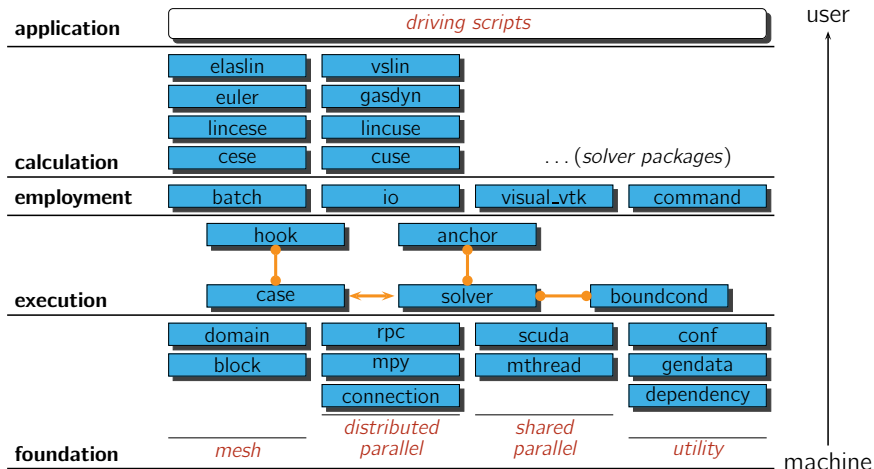
- A standalone class hierarchy (Case) is designed to host the **temporal loop**.




- Hook and Anchor are call-back objects for Case and Solver, respectively.
  - Supplement of main algorithms.
  - Lazy initialization.
  - Facilitating parallel computing and in-situ analysis.



# Overall Design of SOLVCON




# Two Types of Parallel Computing



- Simultaneously use shared-memory and distributed-memory parallel computing (DMPC & DMPC, respectively).
  - Main difference: **Addressing space**.
- Inter-process communication is needed.
  - DMPC is much more complex than SMPC.
  - **DMPC determines the scalability**.
  - **MapReduce is unsuitable**.


# Extending Solver Kernel for SMPC

- A Solver class can be extended to use shared-memory parallel computing.
- Only the spatial loops are modified.
- Can use pthread, OpenMP, CUDA, OpenCL, etc.



# Domain Decomposition for DMPC

- Before computation: Domain decomposition.
  - Use connectivity data to build the graph of cells.
  - Partition the graph by calling METIS or SCOTCH library.
  - Use the partitioned graph to decompose mesh data.

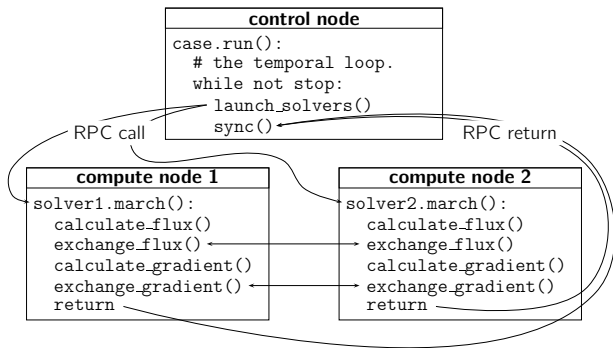


- During computation: Exchange data of the cells on the interface of different sub-domains.
  - Use MPI to communicate among sub-domains.

# Solver Kernels Need Not Know DMPC

- DMPC is in SOLVCON framework.
- SMPC is in solver kernels.

## DMPC Execution Flow in SOLVCON



- When developing solver kernels, we do not need to worry about the complexity of DMPC.
- Hybrid parallelism is achieved by the segregation.

# Post-Processing is Bottleneck

- High-resolution simulations generate a lot of data:
  - For 50 million element mesh, the data for one scalar (single-precision) are **200 MB**.
  - A typical run has at least 10,000 time steps.
  - Transient analysis:  $10,000 \times 200 \text{ MB} = \mathbf{2 \text{ TB}}$ .
  - $\rho, p, T, \vec{v}, \vec{\omega}$  for CFD:  $2 \text{ TB} \times 9 = \mathbf{18 \text{ TB}}$ .
- Workaround: Reducing output frequency.
  - Every 100 time steps: **160 GB**.
- Post-processing the solutions is painfully time-consuming:
  - The large data are usually processed by using a single workstation.
  - Turnaround time could be in months.

# Solutions in SOLVCON

- Parallel I/O.
  - Each sub-domain outputs its own solutions.
  - It is used with **parallel post-processing**.
- In situ visualization.
  - Visualization is being done on the fly with the simulation.
  - Everything happens in memory.
  - Output only graphic files, which are much smaller than the full solution field.
- Parallel I/O and in situ visualization are complementary to each other.

# Driving Scripts

- Driving scripts manage simulations for SOLVCON.
  - A driving script must create a Case object and call its (i) `init()`, (ii) `run()`, and (iii) `cleanup()` methods.
  - No input file is needed.
- Driving scripts can specify **logic** to the simulations in addition to parameters.
  - Anything higher than the foundation layer (the lowest layer) can be **replaced** by code written in driving scripts.
  - Including but not limited to Case, Solver, BC classes, Hook, and Anchor classes.



# Three-Dimensional Sod's Shock Tube

- `$SCR00T/examples/gasdyn/tube`.
  - The one-dimensional version is a standard test case for gas-dynamics codes.
  - This version demonstrates the three-dimensional capability of SOLVCON.
  - You also need CUBIT to generate the mesh.
- Let's demo:
  - `cd $SCR00T/examples/gasdyn/tube`
  - `./go`

# General Structure of a Driving Script

- The driving script for the shock tube problem:

```

1  from solvcon.kerpak import gasdyn
2  class DiaphragmAnchor(gasdyn.GasdynAnchor):
3      ...
4  def tube_base(casename=None,
5               gamma=None, rho1=None, p1=None, rho2=None, p2=None,
6               psteps=None, ssteps=None, **kw
7  ):
8      ...
9  if __name__ == '__main__':
10     cse = tube_base('tube_20', use_incenter=True,
11                   gamma=1.4, rho1=1.0, p1=1.0, rho2=0.125, p2=0.25,
12                   time_increment=1.8e-3, steps_run=400, ssteps=40, psteps=1)
13     cse.init()
14     cse.run()
15     cse.cleanup()

```

- There's another “delayed” mode for driving scripts. See `examples/gasdyn/blnt`.

# Use Anchor for Initial Condition

- The `provide()` method is invoked before the temporal loop:

```
1  class DiaphragmAnchor(gasdyn.GasdynAnchor):  
2      ...  
3      def provide(self):  
4          super(DiaphragmAnchor, self).provide()  
5          gamma = self.gamma  
6          svr = self.svr  
7          svr.soln[:,0].fill(self.rho1)  
8          svr.soln[:,1].fill(0.0)  
9          svr.soln[:,2].fill(0.0)  
10         if svr.ndim == 3:  
11             svr.soln[:,3].fill(0.0)  
12         svr.soln[:,svr.ndim+1].fill(self.p1/(gamma-1))  
13         # set.  
14         slct = svr.clcnd[:,0] > 0.0  
15         svr.soln[slct,0] = self.rho2  
16         svr.soln[slct,svr.ndim+1] = self.p2  
17         # update.  
18         svr.sol[:] = svr.soln[:]
```

# Setup the Flow of the Simulation

- A Case object is created for the simulation:

```
1  def tube_base(...):
2      ...
3      # set up case.
4      basedir = os.path.abspath(os.path.join(os.getcwd(), 'result'))
5      cse = gasdyn.GasdynCase(basedir=basedir, rootdir=env.projdir,
6                             basefn=casename, mesher=mesher, bmap=bcmap, **kw)
7      # informative.
8      cse.runhooks.append(hook.BlockInfoHook)
9      ...
10     # initializer.
11     ...
12     cse.runhooks.append(DiaphragmAnchor,
13                        gamma=gamma, rho1=rho1, p1=p1, rho2=rho2, p2=p2)
14     # post processing.
15     ...
16     cse.runhooks.append(gasdyn.GasdynAnchor, rsteps=ssteps)
17     ...
18     return cse
```

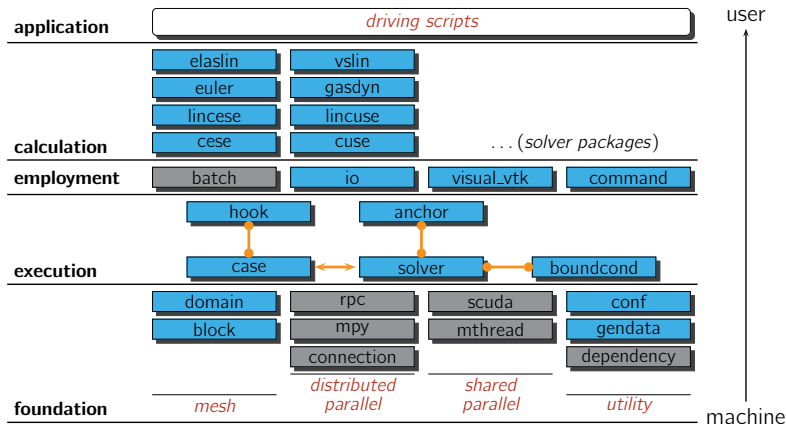
# The Results

- Run the simulation case:

```
1  if __name__ == '__main__':  
2      cse = tube_base('tube_20', use_incenter=True,  
3          gamma=1.4, rho1=1.0, p1=1.0, rho2=0.125, p2=0.25,  
4          time_increment=1.8e-3, steps_run=400, ssteps=40, psteps=1)  
5      cse.init()  
6      cse.run()  
7      cse.cleanup()
```

- Let's launch ParaView to visualize and render the calculated results.

# A Planned Renovation of SOLVCON



- Simply the architecture: Rely more on mpi4py, OpenMP, etc.
- Use Cython instead of ctypes for maintainability.

# Conclusions

- We've used Python to construct a high-performance system for solving conservation laws or PDEs.
  - It's multi-physics: Fluid mechanics and solid mechanics (and electromagnetics in the future).
  - It has a clear architecture.
  - Parallel computing is made easy.
- It's very **productive**.
- And ...

# Python Rocks

Yes, it is!



# Thank You